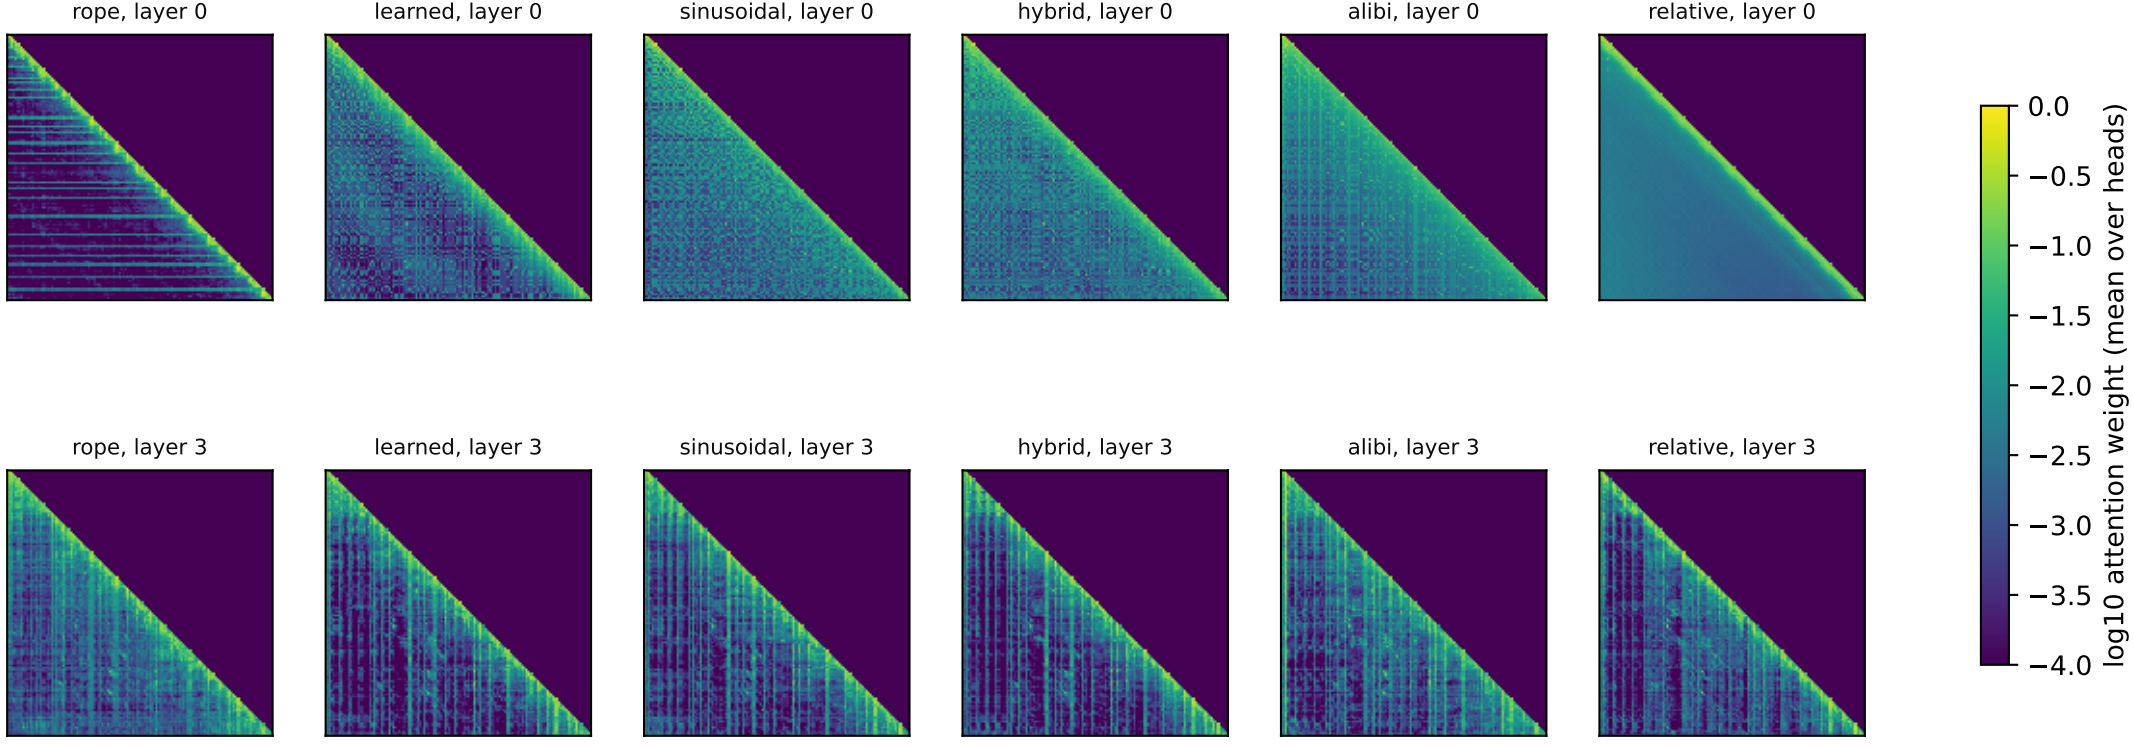

Attention on one test window (W=128, same text for every encoding); rows: query, columns: key
rope, layer 0
learned, layer 0
sinusoidal, layer 0
hybrid, layer 0
alibi, layer 0
relative, layer 0
rope, layer 3
learned, layer 3
sinusoidal, layer 3
hybrid, layer 3
alibi, layer 3
relative, layer 3
0.0
−0.5
−1.0
−1.5
−2.0
−2.5
−3.0
−3.5
−4.0
log10 attention weight (mean over heads)