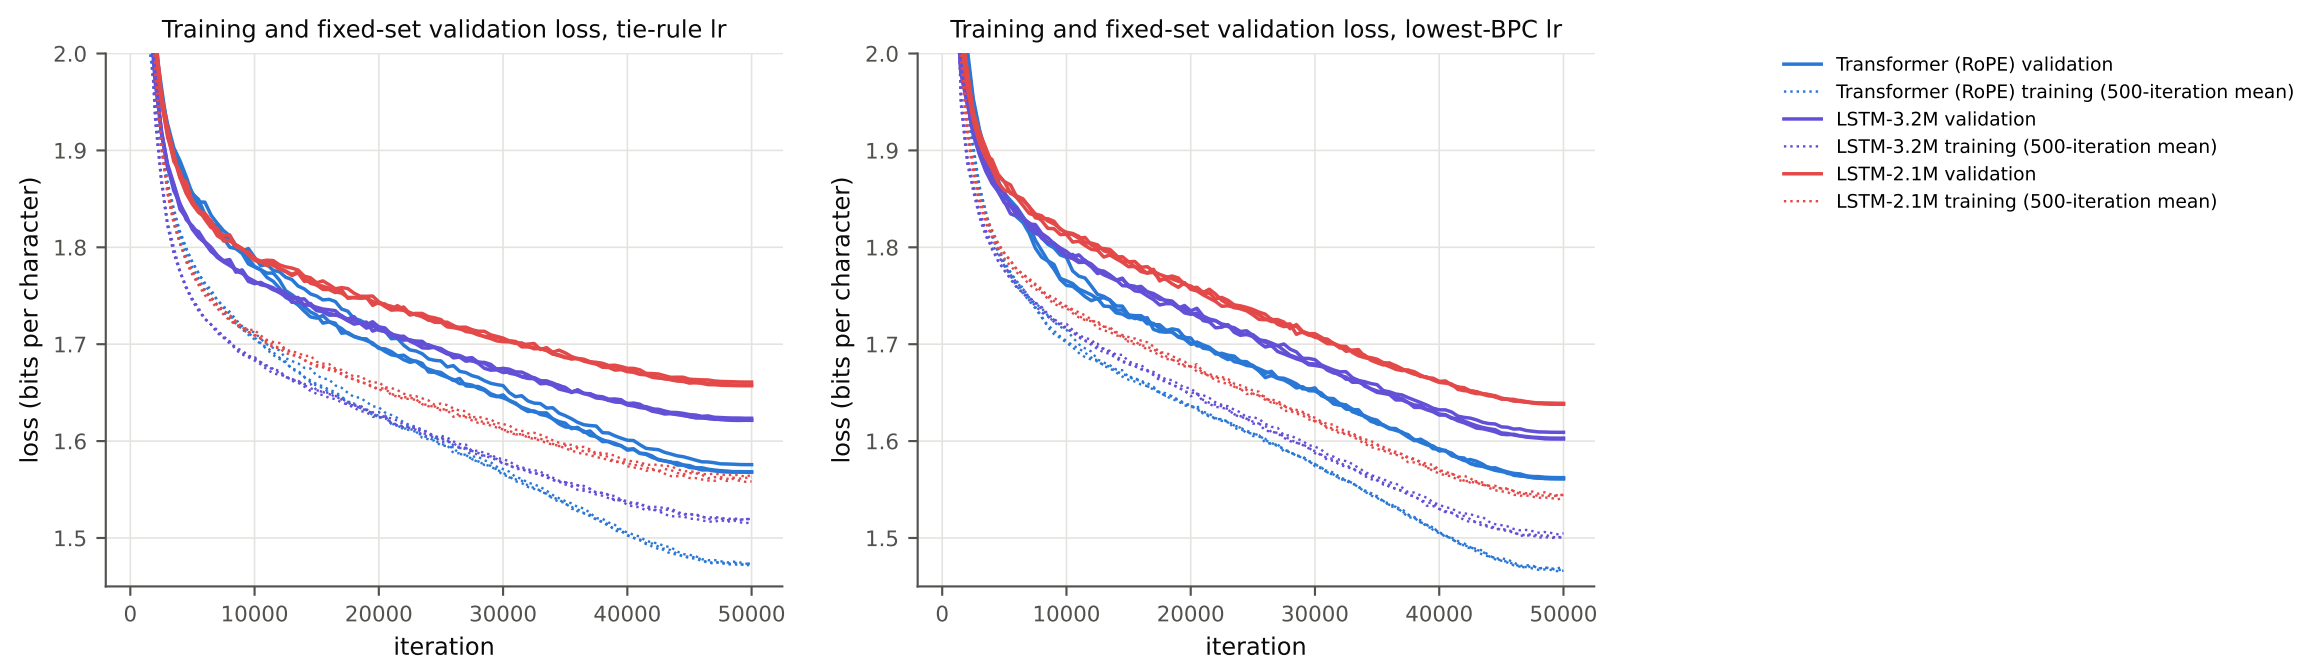
Training and fixed-set validation loss, tie-rule lr
Training and fixed-set validation loss, lowest-BPC lr
loss (bits per character)
iteration
2.0
1.9
1.8
1.7
1.6
1.5
0
10000
20000
30000
40000
50000
Transformer (RoPE) validation
Transformer (RoPE) training (500-iteration mean)
LSTM-3.2M validation
LSTM-3.2M training (500-iteration mean)
LSTM-2.1M validation
LSTM-2.1M training (500-iteration mean)